

Cours 8IAR125: Intelligence Artificielle pour le Jeu Vidéo

TP #1 : Machine à états finis

LE TRAVAIL DOIT SE FAIRE EN EQUIPE DE TROIS (TRINÔME)

À remettre le 27 Septembre 2017¹

1. But

Se familiariser avec les machines à états finis en utilisant :

- Langage C++;
- Un modèle de machine d'états d'un agent-pnj;
- les Designs patterns : State et Singleton;
- La communication de haut niveau à base de messages.

2. Travail à effectuer

Étendre l'application de jeu «chercheur d'OR dans l'ouest», de l'exemple du chapitre 2 du livre de *Mat Buckland*. Plus précisément, de la version «**WestWorldWithMessaging**». Le but consiste à introduire un autre agent-pnj, de votre choix, dans le monde du jeu composé déjà des agents, Bob le chercheur d'or et sa femme Elsa. Les comportements de ces deux agents sont spécifiés par leurs machines d'états respectives, dont l'implémentation est bien documentée dans le livre. Le code source est accessible sur le site http://www.ai-junkie.com/books/toc_pgaibe.html Voir aussi le chapitre qui est annexé² à la page 4.

¹ Dans un des comptes du trinôme en indiquant dans un fichier « readme » les noms et les codes permanents des autres coéquipiers. La remise doit se faire avant minuit (00h00).

² En attendant l'achat du livre, le chapitre 2 est aussi disponible sur le site de l'auteur du livre <http://www.ai-junkie.com/ai-junkie.html>, c'est exceptionnel. Pour les prochains travaux, il n'y aura pas de photocopies. Hâtez vous !

3. Cadre de l'étude

Il s'agit d'introduire un scénario d'une querelle à l'intérieur d'un bar (bistrot) entre l'agent Bob, lorsqu'il est dans un état de soif, et un autre agent ivrogne, par exemple, qui déclenche une altercation lorsqu'il est dans un état de soûlerie. Le scénario doit être simple. Le but est de rajouter à l'application Chapitre2\ **WestWorldWithMessaging**, une nouvelle machine à états-finis qui décrit le comportement du nouveau agent-soulard et par conséquent, de modifier celle de Bob. Ces deux agents doivent communiquer directement par envois de messages en se basant sur la technique déjà implantée. Vous devez donc fournir :

- 3.1 Un schéma du diagramme d'état-transition de la nouvelle machine d'états de l'agent Bob. Vous pouvez utiliser l'outil comme Visio de MS-Office, ou un autre logiciel. (1 point)
- 3.2 Un schéma du diagramme d'états-transition de l'agent-soulard (1 point)
- 3.3 Le code source lisible et l'exécutable de la nouvelle application (8 points)
- 3.4 Quelques captures d'écran (snapshot) de résultats d'exécution de votre programme
- 3.5 L'ajout des fonctionnalités suivantes est facultatif mais ça sera bonifié:
 - 3.5.1 Une interface graphique pour :
 - spécifier l'état de départ d'un agent (2 points);
 - visualiser les messages échangés entre les agents (2 points).
 - 3.5.2 Une implantation d'une architecture multi-threads des agents (10 points).

(NB) Vous devez indiquer dans un fichier texte, toute autre fonctionnalité que vous avez rajoutée à la version fournie. Ce fichier doit être dans votre répertoire de remise de votre travail. **Aucun travail ne sera accepté par un envoi de courriel !**

Annexe sur les directives de remise des travaux

- Le travail doit compiler en VS 2010 au minimum
- Effacez les fichiers *.ncb *.user *.suo et les répertoires Debug et Release en entier.

- Incluez tout ce qui est nécessaire pour que ça compile du premier coup.
- Utilisez des chemins relatifs dans les configurations du projet tel que : « ..\common » .
- Incluez un readme.txt avec les noms des membres de l'équipe ainsi, leurs codes permanents et les problèmes que vous n'avez pas réussis à résoudre.
- Si quelque chose ne marche pas et que vous l'avez bien documenté vous allez perdre moins de points.
- Le travail doit compiler en Debug et en Release donc configurez le projet en conséquence.
- Comprimez votre travail en fichier .ZIP.
- Nommez le travail de la façon suivante TP#-IDENTITÉDUCOMPTE (tel TP1-ROBJ1205.zip) dans lequel le travail a été placé.
- Déposez vos travaux compressés dans votre compte sur la racine.
- Pour les diagrammes fait par Visio, il faudrait les exporter absolument en image PNG compressée

Annexe sur le chapitre 2 du livre disponible sur le site de l'auteur

<http://www.ai-junkie.com/ai-junkie.html>

Finite state machines, or FSMs as they are usually referred to, have for many years been the AI coder's instrument of choice to imbue a game agent with the illusion of intelligence. You will find FSMs of one kind or another in just about every game to hit the shelves since the early days of video games, and despite the increasing popularity of more esoteric agent architectures, they are going to be around for a long time to come. Here are just some of the reasons why:

- **They are quick and simple to code.** There are many ways of programming a finite state machine and almost all of them are reasonably simple to implement. You'll see several alternatives described in this article together with the pros and cons of using *Athem*.
- **They are easy to debug.** Because a game agent's behavior is broken down into easily manageable chunks, if an agent starts acting strangely, it can be debugged by adding tracer code to each state. In this way, the AI programmer can easily follow the sequence of events that precedes the buggy behavior and take action accordingly.
- **They have little computational overhead.** Finite state machines use hardly any precious processor time because they essentially follow hard-coded rules. There is no real "thinking" involved beyond the if-*this*-then-*that* sort of thought process.
- **They are intuitive.** It's human nature to think about things as being in one state or another and we often refer to ourselves as being in such and such a state. How many times have you "got yourself into a state" or found yourself in "the right state of mind"? Humans don't really work like finite state machines of course, but sometimes we find it useful to think of our behavior in this way. Similarly, it is fairly easy to break down a game agent's behavior into a number of states and to create the rules required for manipulating them. For the same reason, finite state machines also make it easy for you to discuss the design of your AI with non-programmers (with game producers and level designers for example), providing improved communication and exchange of ideas.
- **They are flexible.** A game agent's finite state machine can easily be adjusted and tweaked by the programmer to provide the behavior required by the game designer. It's also a simple matter to expand the scope of an agent's behavior by adding new states and rules. In addition, as your AI skills grow you'll find that finite state machines provide a solid backbone with which you can combine other techniques such as fuzzy logic or neural networks.

What Exactly Is a Finite State Machine?

Historically, a finite state machine is a rigidly formalized device used by mathematicians to solve problems. The most famous finite state machine is probably Alan Turing's hypothetical device: the Turing machine, which he wrote about in his 1936 paper, "On Computable Numbers." This was a machine presaging modern-day programmable computers that could perform any logical operation by reading, writing, and erasing symbols on an infinitely long strip of tape. Fortunately, as AI programmers, we can forgo the formal mathematical definition of a finite state machine; a descriptive one will suffice:

A finite state machine is a device, or a model of a device, which has a finite number of states it can be in at any given time and can operate on input to either make transitions from one state to another or to cause an output or action to take place. A finite state machine can only be in one state at any moment in time.

The idea behind a finite state machine, therefore, is to decompose an object's behavior into easily manageable "chunks" or states. The light switch on your wall, for example, is a very simple finite state machine. It has two states: on and off. Transitions between states are made by the input of your finger. By flicking the switch up it makes the transition from off to on, and by flicking the switch down it makes the transition from on to off. There is no output or action associated with the off state (unless you consider the bulb being off as an action), but when it is in the on state electricity is allowed to flow through the switch and light up your room via the filament in a lightbulb. See Figure 2.1.




Figure 2.1. A light switch is a finite state machine. (Note that the switches are reversed in Europe and many other parts of the world.)

Of course, the behavior of a game agent is usually much more complex than a lightbulb (thank goodness!). Here are some examples of how finite state machines have been used in games.

- The ghosts' behavior in Pac-Man is implemented as a finite state machine. There is one Evade state, which is the same for all ghosts, and then each ghost has its own Chase state, the actions of which are implemented differently for each ghost. The input of the player eating one of the power pills is the condition for the transition from Chase to Evade. The input of a timer running down is the condition for the transition from Evade to Chase.
- Quake-style bots are implemented as finite state machines. They have states such as FindArmor, FindHealth, SeekCover, and RunAway. Even the weapons in Quake implement their own mini finite state machines. For example, a rocket may implement states such as Move, TouchObject, and Die.
- Players in sports simulations such as the soccer game FIFA2002 are implemented as state machines. They have states such as Strike, Dribble, ChaseBall, and MarkPlayer. In addition, the teams themselves are often implemented as FSMs and can have states such as KickOff, Defend, or WalkOutOnField.
- The NPCs (non-player characters) in RTSs (real-time strategy games) such as Warcraft make use of finite state machines. They have states such as MoveToPosition, Patrol, and FollowPath.

Implementing a Finite State Machine

There are a number of ways of implementing finite state machines. A naive approach is to use a series of if-then statements or the slightly tidier mechanism of a switch statement. Using a switch with an enumerated type to represent the states looks something like this:

```
enum StateType{state_RunAway, state_Patrol, state_Attack};

void Agent::UpdateState(StateType CurrentState)
{
    switch(CurrentState)
    {
        case state_RunAway:

            EvadeEnemy();

            if (Safe())
            {
                ChangeState(state_Patrol);
            }

            break;

        case state_Patrol:
```

```

FollowPatrolPath();

if (Threatened())
{
    if (StrongerThanEnemy())
    {
        ChangeState(state_Attack);
    }
    else
    {
        ChangeState(state_RunAway);
    }
}

break;

case state_Attack:

    if (WeakerThanEnemy())
    {
        ChangeState(state_RunAway);
    }
    else
    {
        BashEnemyOverHead();
    }

    break;

} //end switch
}

```

Although at first glance this approach seems reasonable, when applied practically to anything more complicated than the simplest of game objects, the switch/if-then solution becomes a monster lurking in the shadows waiting to pounce. As more states and conditions are added, this sort of structure ends up looking like spaghetti very quickly, making the program flow difficult to understand and creating a debugging nightmare. In addition, it's inflexible and difficult to extend beyond the scope of its original design, should that be desirable... and as we all know, it most often is. Unless you are designing a state machine to implement very simple behavior (or you are a genius), you will almost certainly find yourself first tweaking the agent to cope with unplanned-for circumstances before honing the behavior to get the results you thought you were going to get when you first planned out the state machine!

Additionally, as an AI coder, you will often require that a state perform a specific action (or actions) when it's initially entered or when the state is exited. For example, when an agent enters the state RunAway you may want it to wave its arms in the air and scream "Arghhhhhhh!" When it finally escapes and changes state to Patrol, you may want it to emit a sigh, wipe its forehead, and say "Phew!" These are actions that only occur when the RunAway state is entered or exited and not during the usual update step. Consequently, this additional functionality must ideally be built into your state machine architecture. To do this within the framework of a switch or if-then

architecture would be accompanied by lots of teeth grinding and waves of nausea, and produce very ugly code indeed.

State Transition Tables

A better mechanism for organizing states and affecting state transitions is a *state transition table*. This is just what it says it is: a table of conditions and the states those conditions lead to. Table 2.1 shows an example of the mapping for the states and conditions shown in the previous example.

Table 2.1. A simple state transition table

Current State	Condition	State Transition
Runaway	Safe	Patrol
Attack	WeakerThanEnemy	RunAway
Patrol	Threatened AND StrongerThanEnemy	Attack
Patrol	Threatened AND WeakerThanEnemy	RunAway

This table can be queried by an agent at regular intervals, enabling it to make any necessary state transitions based on the stimulus it receives from the game environment. Each state can be modeled as a separate object or function existing external to the agent, providing a clean and flexible architecture. One that is much less prone to spaghettiification than the if-then/switch approach discussed in the previous section.

Someone once told me a vivid and silly visualization can help people to understand an abstract concept. Let's see if it works...

Imagine a robot kitten. It's shiny yet cute, has wire for whiskers and a slot in its stomach where cartridges — analogous to its states — can be plugged in. Each of these cartridges is programmed with logic, enabling the kitten to perform a specific set of actions. Each set of actions encodes a different behavior; for example, `play_with_string`, `eat_fish`, or `poo_on_carpet`. Without a cartridge stuffed inside its belly the kitten is an inanimate metallic sculpture, only able to sit there and look cute... in a Metal Mickey kind of way.

The kitten is very dexterous and has the ability to autonomously exchange its cartridge for another if instructed to do so. By providing the rules that dictate when a cartridge should be switched, it's possible to string together sequences of cartridge insertions permitting the creation of all sorts of interesting and complicated behavior. These rules are programmed onto a tiny chip situated inside the kitten's head, which is analogous to the state transition table we discussed earlier. The chip communicates with the kitten's internal functions to retrieve the information necessary to process the rules (such as how hungry Kitty is or how playful it's feeling). As a result, the state transition chip can be programmed with rules like:

IF `Kitty_Hungry` AND NOT `Kitty_Playful` SWITCH_CARTRIDGE `eat_fish`

All the rules in the table are tested each time step and instructions are sent to Kitty to switch cartridges accordingly.

This type of architecture is very flexible, making it easy to expand the kitten's repertoire by adding new cartridges. Each time a new cartridge is added, the owner is only required to take a

screwdriver to the kitten's head in order to remove and reprogram the state transition rule chip. It is not necessary to interfere with any other internal circuitry.

Embedded Rules

An alternative approach is to *embed the rules for the state transitions within the states themselves*. Applying this concept to Robo-Kitty, the state transition chip can be dispensed with and the rules moved directly into the cartridges. For instance, the cartridge for `play_with_string` can monitor the kitty's level of hunger and instruct it to switch cartridges for the `eat_fish` cartridge when it senses hunger rising. In turn the `eat_fish` cartridge can monitor the kitten's bowel and instruct it to switch to the `poo_on_carpet` cartridge when it senses poo levels are running dangerously high.

Although each cartridge may be aware of the existence of any of the other cartridges, each is a self-contained unit and not reliant on any external logic to decide whether or not it should allow itself to be swapped for an alternative. As a consequence, it's a straightforward matter to add states or even to swap the whole set of cartridges for a completely new set (maybe ones that make little Kitty behave like a raptor). There's no need to take a screwdriver to the kitten's head, only to a few of the cartridges themselves.

Let's take a look at how this approach is implemented within the context of a video game. Just like Kitty's cartridges, states are encapsulated as objects and contain the logic required to facilitate state transitions. In addition, all state objects share a common interface: a pure virtual class named `State`. Here's a version that provides a simple interface:

```
class State
{
public:

    virtual void Execute (Troll* troll) = 0;
};
```

Now imagine a `Troll` class that has member variables for attributes such as health, anger, stamina, etc., and an interface allowing a client to query and adjust those values. A `Troll` can be given the functionality of a finite state machine by adding a pointer to an instance of a derived object of the `State` class, and a method permitting a client to change the instance the pointer is pointing to.

```
class Troll
{
    /* ATTRIBUTES OMITTED */

    State* m_pCurrentState;

public:

    /* INTERFACE TO ATTRIBUTES OMITTED */

    void Update()
    {
        m_pCurrentState->Execute(this);
    }
}
```

```

void ChangeState(const State* pNewState)
{
    delete m_pCurrentState;
    m_pCurrentState = pNewState;
}
};

```

When the Update method of a Troll is called, it in turn calls the Execute method of the current state type with the this pointer. The current state may then use the Troll interface to query its owner, to adjust its owner's attributes, or to effect a state transition. In other words, how a Troll behaves when updated can be made completely dependent on the logic in its current state. This is best illustrated with an example, so let's create a couple of states to enable a troll to run away from enemies when it feels threatened and to sleep when it feels safe.

```

//-----State_Runaway
class State_RunAway : public State
{
public:

    void Execute(Troll* troll)
    {
        if (troll->isSafe())
        {
            troll->ChangeState(new State_Sleep());
        }
        else
        {
            troll->MoveAwayFromEnemy();
        }
    }
};

//-----State_Sleep
class State_Sleep : public State
{
public:

    void Execute(Troll* troll)
    {
        if (troll->isThreatened())
        {
            troll->ChangeState(new State_RunAway())
        }

        else
        {
            troll->Snore();
        }
    }
}

```

```
};
```

As you can see, when updated, a troll will behave differently depending on which of the states `m_pCurrentState` points to. Both states are encapsulated as objects and both provide the rules effecting state transition. All very neat and tidy.

This architecture is known as the *state design pattern* and provides an elegant way of implementing state-driven behavior. Although this is a departure from the mathematical formalization of an FSM, it is intuitive, simple to code, and easily extensible. It also makes it extremely easy to add enter and exit actions to each state; all you have to do is create Enter and Exit methods and adjust the agent's `ChangeState` method accordingly. You'll see the code that does exactly this very shortly.

The West World Project

As a practical example of how to create agents that utilize finite state machines, we are going to look at a game environment where agents inhabit an Old West-style gold mining town named West World. Initially there will only be one inhabitant — a gold miner named Miner Bob — but later his wife will also make an appearance. You will have to imagine the tumbleweeds, creakin' mine props, and desert dust blowin' in your eyes because West World is implemented as a simple text-based console application. Any state changes or output from state actions will be sent as text to the console window. I'm using this plaintext-only approach as it demonstrates clearly the mechanism of a finite state machine without adding the code clutter of a more complex environment.

There are four locations in West World: a *goldmine*, a *bank* where Bob can deposit any nuggets he finds, a *saloon* in which he can quench his thirst, and *home-sweet-home* where he can sleep the fatigue of the day away. Exactly where he goes, and what he does when he gets there, is determined by Bob's current state. He will change states depending on variables like thirst, fatigue, and how much gold he has found hacking away down in the gold mine.

Before we delve into the source code, check out the following sample output from the `WestWorld1` executable.

```
Miner Bob: Pickin' up a nugget
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Goin' to the bank. Yes siree
Miner Bob: Depositin' gold. Total savings now: 3
Miner Bob: Leavin' the bank
Miner Bob: Walkin' to the gold mine
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
```

```
Miner Bob: Boy, ah sure is thusty! Walkin' to the saloon
Miner Bob: That's mighty fine sippin' liquor
Miner Bob: Leavin' the saloon, feelin' good
Miner Bob: Walkin' to the gold mine
Miner Bob: Pickin' up a nugget
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Goin' to the bank. Yes siree
Miner Bob: Depositin' gold. Total savings now: 4
Miner Bob: Leavin' the bank
Miner Bob: Walkin' to the gold mine
Miner Bob: Pickin' up a nugget
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Boy, ah sure is thusty! Walkin' to the saloon
Miner Bob: That's mighty fine sippin' liquor
Miner Bob: Leavin' the saloon, feelin' good
Miner Bob: Walkin' to the gold mine
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Goin' to the bank. Yes siree
Miner Bob: Depositin' gold. Total savings now: 5
Miner Bob: Woohoo! Rich enough for now. Back home to mah li'l lady
Miner Bob: Leavin' the bank
Miner Bob: Walkin' home
Miner Bob: ZZZZ...
Miner Bob: ZZZZ...
Miner Bob: ZZZZ...
Miner Bob: ZZZZ...
Miner Bob: What a God-darn fantastic nap! Time to find more gold
```

In the output from the program, each time you see Miner Bob change location he is changing state. All the other events are the actions that take place within the states. We'll examine each of Miner Bob's potential states in just a moment, but for now, let me explain a little about the code structure of the demo.

(You can download the accompanying project files [from my site](#))

The BaseGameEntity Class

All inhabitants of West World are derived from the base class BaseGameEntity. This is a simple class with a private member for storing an ID number. It also specifies a pure virtual member function, Update, which must be implemented by all subclasses. Update is a function that gets called every update step and will be used by subclasses to update their state machine along with any other data that must be updated each time step.

The BaseGameEntity class declaration looks like this:

```
class BaseGameEntity
{
private:

    //every entity has a unique identifying number
    int      m_ID;

    //this is the next valid ID. Each time a BaseGameEntity is instantiated
    //this value is updated
    static int m_iNextValidID;

    //this is called within the constructor to make sure the ID is set
    //correctly. It verifies that the value passed to the method is greater
    //or equal to the next valid ID, before setting the ID and incrementing
    //the next valid ID
    void SetID(int val);

public:

    BaseGameEntity(int id)
    {
        SetID(id);
    }

    virtual ~BaseGameEntity(){}

    //all entities must implement an update function
    virtual void Update()=0;

    int      ID()const{return m_ID;}
};
```

For reasons that will become obvious later [in the book], it's very important for each entity in your game to have a unique identifier. Therefore, on instantiation, the ID passed to the constructor is tested in the SetID method to make sure it's unique. If it is not, the program will exit with an assertion failure. In the example given, the entities will use an enumerated value as their unique identifier. These can be found in the file EntityNames.h as ent_Miner_Bob and ent_Elsa.

The Miner Class

The Miner class is derived from the BaseGameEntity class and contains data members representing the various attributes a Miner possesses, such as its health, its level of fatigue, its position, and so forth. Like the troll example shown earlier, a Miner owns a pointer to an instance of a State class in addition to a method for changing what State that pointer points to.

```
class Miner : public BaseGameEntity
{
private:

    //a pointer to an instance of a State
    State*      m_pCurrentState;

    // the place where the miner is currently situated
    location_type    m_Location;

    //how many nuggets the miner has in his pockets
    int             m_iGoldCarried;

    //how much money the miner has deposited in the bank
    int             m_iMoneyInBank;

    //the higher the value, the thirstier the miner
    int             m_iThirst;

    //the higher the value, the more tired the miner
    int             m_iFatigue;

public:

    Miner(int ID);

    //this must be implemented
    void Update();

    //this method changes the current state to the new state
    void ChangeState(State* pNewState);

    /* bulk of interface omitted */
};
```

The Miner::Update method is straightforward; it simply increments the m_iThirst value before calling the Execute method of the current state. It looks like this:

```
void Miner::Update()
{
    m_iThirst += 1;

    if (m_pCurrentState)
    {
```

```
m_pCurrentState->Execute(this);  
}  
}
```

Now that you've seen how the Miner class operates, let's take a look at each of the states a miner can find itself in.

The Miner States

The gold miner will be able to enter one of four states. Here are the names of those states followed by a description of the actions and state transitions that occur within those states:

- **EnterMineAndDigForNugget:** If the miner is not located at the gold mine, he changes location. If already at the gold mine, he digs for nuggets of gold. When his pockets are full, Bob changes state to **VisitBankAndDepositGold**, and if while digging he finds himself thirsty, he will stop and change state to **QuenchThirst**.
- **VisitBankAndDepositGold:** In this state the miner will walk to the bank and deposit any nuggets he is carrying. If he then considers himself wealthy enough, he will change state to **GoHomeAndSleepTilRested**. Otherwise he will change state to **EnterMineAndDigForNugget**.
- **GoHomeAndSleepTilRested:** In this state the miner will return to his shack and sleep until his fatigue level drops below an acceptable level. He will then change state to **EnterMineAndDigForNugget**.
- **QuenchThirst:** If at any time the miner feels thirsty (diggin' for gold is thursty work, don't ya know), he changes to this state and visits the saloon in order to buy a whiskey. When his thirst is quenched, he changes state to **EnterMineAndDigForNugget**.

Sometimes it's hard to follow the flow of the state logic from reading a text description like this, so it's often helpful to pick up pen and paper and draw a *state transition diagram* for your game agents. Figure 2.2 shows the state transition diagram for the gold miner. The bubbles represent the individual states and the lines between them the available transitions.

A diagram like this is better on the eyes and can make it much easier to spot any errors in the logic flow.

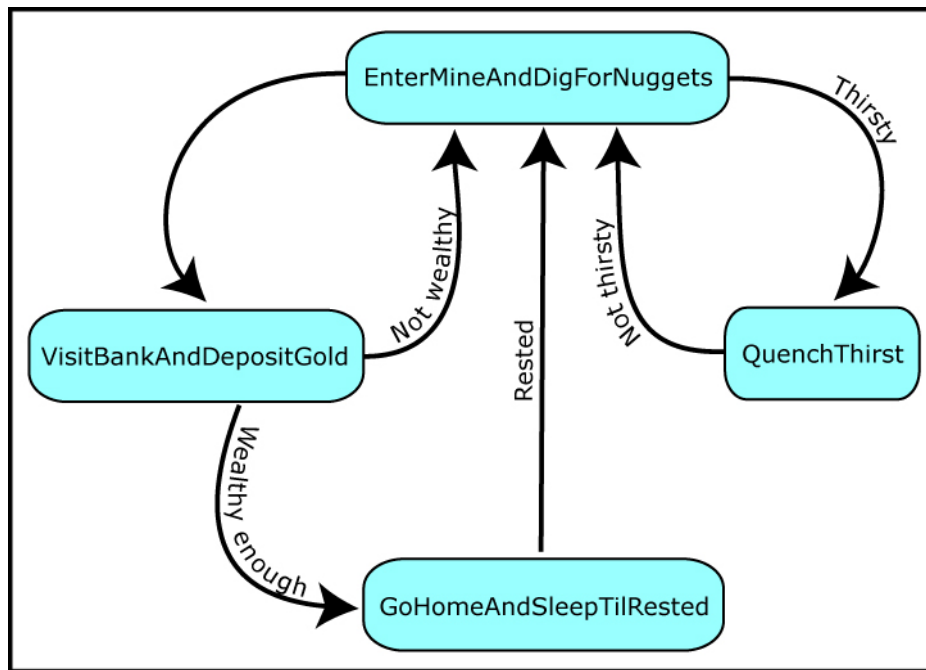


Figure 2.2. Miner Bob's state transition diagram

The State Design Pattern Revisited

You saw a brief description of this design pattern earlier, but it won't hurt to recap. Each of a game agent's states is implemented as a unique class and each agent holds a pointer to an instance of its current state. An agent also implements a *ChangeState* member function that can be called to facilitate the switching of states whenever a state transition is required. The logic for determining any state transitions is contained within each State class. All state classes are derived from an abstract base class, thereby defining a common interface. So far so good. You know this much already.

Earlier it was mentioned that it's usually favorable for each state to have associated Enter and Exit actions. This permits the programmer to write logic that is only executed once at state entry or exit and increases the flexibility of an FSM a great deal. With these features in mind, let's take a look at an enhanced State base class.

```

class State
{
public:

    virtual ~State(){}

    //this will execute when the state is entered
    virtual void Enter(Miner*)=0;

    //this is called by the miner's update function each update-step
    virtual void Execute(Miner*)=0;

    //this will execute when the state is exited
    virtual void Exit(Miner*)=0;
}

```


These additional methods are only called when a Miner changes state. When a state transition occurs, the Miner::ChangeState method first calls the Exit method of the current state, then it assigns the new state to the current state, and finishes by calling the Enter method of the new state (which is now the current state). I think code is clearer than words in this instance, so here's the listing for the ChangeState method:

```
void Miner::ChangeState(State* pNewState)
{
    //make sure both states are valid before attempting to
    //call their methods
    assert (m_pCurrentState && pNewState);

    //call the exit method of the existing state
    m_pCurrentState->Exit(this);

    //change state to the new state
    m_pCurrentState = pNewState;

    //call the entry method of the new state
    m_pCurrentState->Enter(this);
}
```

Notice how a Miner passes the this pointer to each state, enabling the state to use the Miner interface to access any relevant data.

TIP: The state design pattern is also useful for structuring the main components of your game flow. For example, you could have a menu state, a save state, a paused state, an options state, a run state, etc.

Each of the four possible states a Miner may access are derived from the State class, giving us these concrete classes: EnterMineAndDigForNugget, VisitBankAndDepositGold, GoHomeAndSleepTilRested, and QuenchThirst. The Miner::m_pCurrentState pointer is able to point to any of these states. When the Update method of Miner is called, it in turn calls the Execute method of the currently active state with the this pointer as a parameter. These class relationships may be easier to understand if you examine the simplified UML class diagram shown in Figure 2.3. (Click [here](#) for an introduction to UML class diagrams)

Each concrete state is implemented as a singleton object. This is to ensure that there is only one instance of each state, which agents share (those of you unsure of what a singleton is, please read [this](#)). Using singletons makes the design more efficient because they remove the need to allocate and deallocate memory every time a state change is made. This is particularly important if you have many agents sharing a complex FSM and/or you are developing for a machine with limited resources.

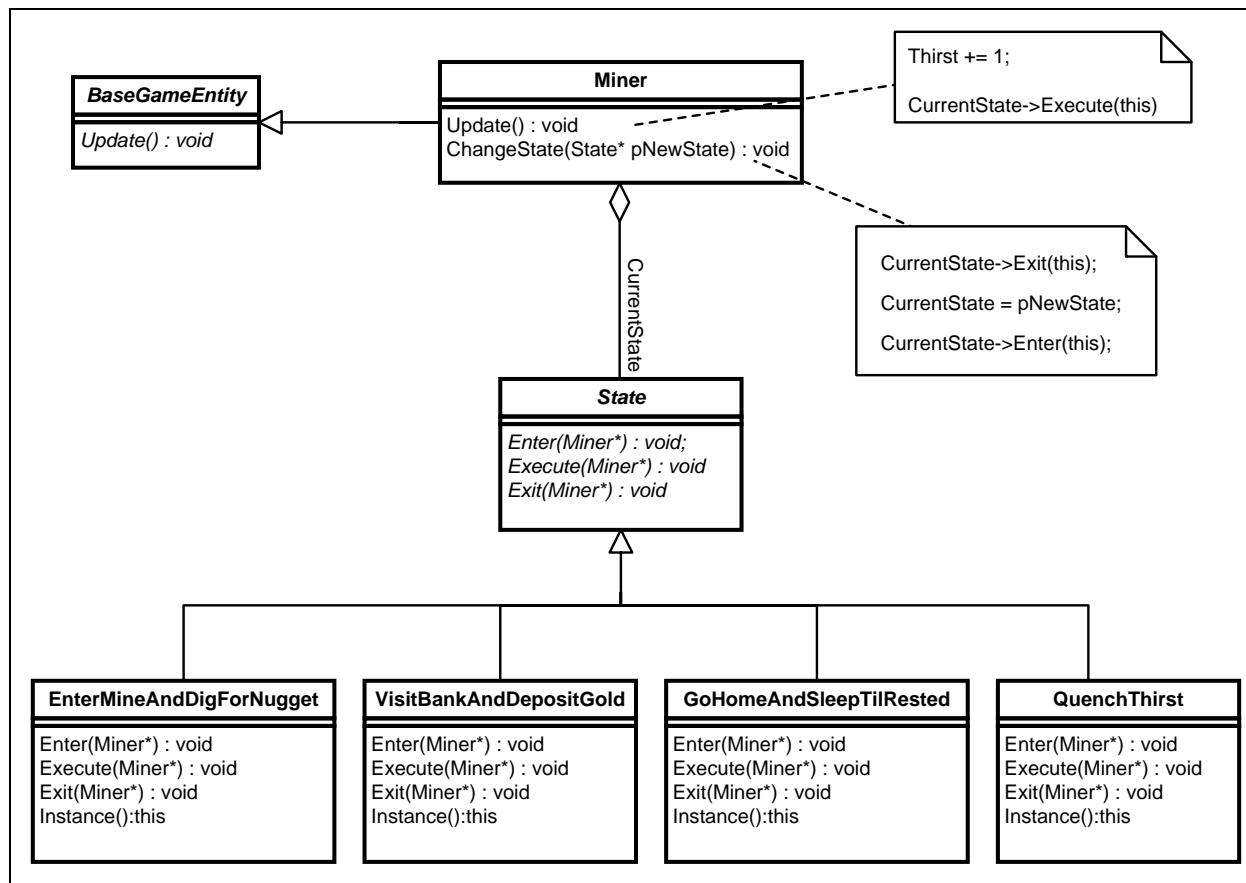


Figure 2.3. UML class diagram for Miner Bob's state machine implementation

NOTE I prefer to use singletons for the states for the reasons I've already given, but there is one drawback. Because they are shared between clients, singleton states are unable to make use of their own local, agent-specific data. For instance, if an agent uses a state that when entered should move it to an arbitrary position, the position cannot be stored in the state itself (because the position may be different for each agent that is using the state). Instead, it would have to be stored somewhere externally and be accessed by the state via the agent's interface. This is not really a problem if your states are accessing only one or two pieces of data, but if you find that the states you have designed are repeatedly accessing lots of external data, it's probably worth considering disposing of the singleton design and writing a few lines of code to manage the allocation and deallocation of state memory.

Okay, let's see how everything fits together by examining the complete code for one of the miner states.

The EnterMineAndDigForNugget State

In this state the miner should change location to be at the gold mine. Once at the gold mine he should dig for gold until his

pockets are full, when he should change state to VisitBankAndDepositNugget. If the miner gets thirsty while digging he should change state to QuenchThirst.

Because concrete states simply implement the interface defined in the virtual base class State, their declarations are very straightforward:

```
class EnterMineAndDigForNugget : public State
{
private:

    EnterMineAndDigForNugget(){}

    /* copy ctor and assignment op omitted */

public:

    //this is a singleton
    static EnterMineAndDigForNugget* Instance();

    virtual void Enter(Miner* pMiner);

    virtual void Execute(Miner* pMiner);

    virtual void Exit(Miner* pMiner);
};
```

As you can see, it's just a formality. Let's take a look at each of the methods in turn.

EnterMineAndDigForNugget::Enter

The code for the Enter method of EnterMineAndDigForNugget is as follows:

```
void EnterMineAndDigForNugget::Enter(Miner* pMiner)
{
    //if the miner is not already located at the goldmine, he must
    //change location to the gold mine
    if (pMiner->Location() != goldmine)
    {
        cout << "\n" << GetNameOfEntity(pMiner->ID()) << ": "
             << "Walkin' to the goldmine";

        pMiner->ChangeLocation(goldmine);
    }
}
```

This method is called when a miner first enters the EnterMineAndDigForNugget state. It ensures that the gold miner is located at the gold mine. An agent stores its location as an enumerated type and the ChangeLocation method changes this value to switch locations.

EnterMineAndDigForNugget::Execute

The Execute method is a little more complicated and contains logic that can change a miner's state. (Don't forget that Execute is the method called each update step from Miner::Update.)

```
void EnterMineAndDigForNugget::Execute(Miner* pMiner)
{
    //the miner digs for gold until he is carrying in excess of MaxNuggets.
    //If he gets thirsty during his digging he stops work and
    //changes state to go to the saloon for a beer.
    pMiner->AddToGoldCarried(1);

    //digging is hard work
    pMiner->IncreaseFatigue();

    cout << "\n" << GetNameOfEntity(pMiner->ID()) << ": "
         << "Pickin' up a nugget";

    //if enough gold mined, go and put it in the bank
    if (pMiner->PocketsFull())
    {
        pMiner->ChangeState(VisitBankAndDepositGold::Instance());
    }

    //if thirsty go and get a beer
    if (pMiner->Thirsty())
    {
        pMiner->ChangeState(QuenchThirst::Instance());
    }
}
```

Note here how the Miner::ChangeState method is called using QuenchThirst's or VisitBankAndDepositGold's Instance member, which provides a pointer to the unique instance of that class.

EnterMineAndDigForNugget::Exit

The Exit method of EnterMineAndDigForNugget outputs a message telling us that the gold miner is leaving the mine.

```
void EnterMineAndDigForNugget::Exit(Miner* pMiner)
{
    cout << "\n" << GetNameOfEntity(pMiner->ID()) << ": "
         << "Ah'm leavin' the goldmine with mah pockets full o' sweet gold";
}
```

I hope an examination of the preceding three methods helps clear up any confusion you may have been experiencing and that you can now see how each state is able to modify the behavior of an agent or effect a transition into another state. You may find it useful at this stage to load up the WestWorld1 project into your IDE and scan the code. In particular, check out all the states in MinerOwnedStates.cpp and examine the Miner class to familiarize yourself with its member variables. Above all else, make sure you understand how the state design pattern works before you

read any further. If you are a little unsure, please take the time to go over the previous text until you feel comfortable with the concept.

You have seen how the use of the state design pattern provides a very flexible mechanism for state-driven agents. It's extremely easy to add additional states as and when required. Indeed, should you so wish, you can switch an agent's entire state architecture for an alternative one. This can be useful if you have a very complicated design that would be better organized as a collection of several separate smaller state machines. For example, the state machine for a first-person shooter (FPS) like Unreal 2 tends to be large and complex. When designing the AI for a game of this sort you may find it preferable to think in terms of several smaller state machines representing functionality like "defend the flag" or "explore map," which can be switched in and out when appropriate. The state design pattern makes this easy to do.

Making the State Base Class Reusable

As the design stands, it's necessary to create a separate State base class for each character type to derive its states from. Instead, let's make it reusable by turning it into a class template.

```
template <class entity_type>
class State
{
public:

    virtual void Enter(entity_type*)=0;

    virtual void Execute(entity_type*)=0;

    virtual void Exit(entity_type*)=0;

    virtual ~State(){}
};
```

The declaration for a concrete state — using the EnterMineAndDigForNugget miner state as an example — now looks like this:

```
class EnterMineAndDigForNugget : public State<Miner>
{
public:

    /* OMITTED */
};
```

This, as you will see shortly, makes life easier in the long run.

Global States and State Blips

More often than not, when designing finite state machines you will end up with code that is duplicated in every state. For example, in the popular game The Sims by Maxis, a Sim may feel the urge of nature come upon it and have to visit the bathroom to relieve itself. This urge may occur in any state the Sim may be in and at any time. Given the current design, to bestow the gold miner

with this type of behavior, duplicate conditional logic would have to be added to every one of his states, or alternatively, placed into the `Miner::Update` function. While the latter solution is acceptable, it's better to create a *global state* that is called every time the FSM is updated. That way, all the logic for the FSM is contained within the states and not in the agent class that owns the FSM.

To implement a global state, an additional member variable is required:

```
//notice how now that State is a class template we have to declare the entity type
State<Miner>* m_pGlobalState;
```

In addition to global behavior, occasionally it will be convenient for an agent to enter a state with the condition that when the state is exited, the agent returns to its previous state. I call this behavior a *state blip*. For example, just as in The Sims, you may insist that your agent can visit the bathroom at any time, yet make sure it always returns to its prior state. To give an FSM this type of functionality it must keep a record of the previous state so the state blip can revert to it. This is easy to do as all that is required is another member variable and some additional logic in the `Miner::ChangeState` method.

By now though, to implement these additions, the `Miner` class has acquired two extra member variables and one additional method. It has ended up looking something like this (extraneous detail omitted):

```
class Miner : public BaseGameEntity
{
private:

    State<Miner>* m_pCurrentState;
    State<Miner>* m_pPreviousState;
    State<Miner>* m_pGlobalState;
    ...

public:

    void ChangeState(State<Miner>* pNewState);
    void RevertToPreviousState();
    ...
};
```

Hmm, looks like it's time to tidy up a little.

Creating a State Machine Class

The design can be made a lot cleaner by encapsulating all the state related data and methods into a state machine class. This way an agent can own an instance of a state machine and delegate the management of current states, global states, and previous states to it.

With this in mind take a look at the following `StateMachine` class template.

```
template <class entity_type>
class StateMachine
{
```

```

private:

    //a pointer to the agent that owns this instance
    entity_type*    m_pOwner;

    State<entity_type>* m_pCurrentState;

    //a record of the last state the agent was in
    State<entity_type>* m_pPreviousState;

    //this state logic is called every time the FSM is updated
    State<entity_type>* m_pGlobalState;

public:

    StateMachine(entity_type* owner):m_pOwner(owner),
                                     m_pCurrentState(NULL),
                                     m_pPreviousState(NULL),
                                     m_pGlobalState(NULL)
    {}

    //use these methods to initialize the FSM
    void SetCurrentState(State<entity_type>* s){m_pCurrentState = s;}
    void SetGlobalState(State<entity_type>* s) {m_pGlobalState = s;}
    void SetPreviousState(State<entity_type>* s){m_pPreviousState = s;}

    //call this to update the FSM
    void Update()const
    {
        //if a global state exists, call its execute method
        if (m_pGlobalState) m_pGlobalState->Execute(m_pOwner);

        //same for the current state
        if (m_pCurrentState) m_pCurrentState->Execute(m_pOwner);
    }

    //change to a new state
    void ChangeState(State<entity_type>* pNewState)
    {
        assert(pNewState &&
               "<StateMachine::ChangeState>: trying to change to a null state");

        //keep a record of the previous state
        m_pPreviousState = m_pCurrentState;

        //call the exit method of the existing state
        m_pCurrentState->Exit(m_pOwner);

        //change state to the new state
        m_pCurrentState = pNewState;
    }

```

```

    //call the entry method of the new state
    m_pCurrentState->Enter(m_pOwner);
}

//change state back to the previous state
void RevertToPreviousState()
{
    ChangeState(m_pPreviousState);
}

//accessors
State<entity_type>* CurrentState() const{return m_pCurrentState;}
State<entity_type>* GlobalState() const{return m_pGlobalState;}
State<entity_type>* PreviousState() const{return m_pPreviousState;}

//returns true if the current state's type is equal to the type of the
//class passed as a parameter.
bool isInState(const State<entity_type>& st)const;
};

```

Now all an agent has to do is to own an instance of a StateMachine and implement a method to update the state machine to get full FSM functionality.

The improved Miner class now looks like this:

```

class Miner : public BaseGameEntity
{
private:

    //an instance of the state machine class
    StateMachine<Miner>* m_pStateMachine;

    /* EXTRANEIOUS DETAIL OMITTED */

public:

    Miner(int id):m_Location(shack),
        m_iGoldCarried(0),
        m_iMoneyInBank(0),
        m_iThirst(0),
        m_iFatigue(0),
        BaseGameEntity(id)

    {
        //set up state machine
        m_pStateMachine = new StateMachine<Miner>(this);

        m_pStateMachine->SetCurrentState(GoHomeAndSleepTilRested::Instance());
        m_pStateMachine->SetGlobalState(MinerGlobalState::Instance());
    }
}

```



```

}

~Miner(){delete m_pStateMachine;}

void Update()
{
    ++m_iThirst;
    m_pStateMachine->Update();
}

StateMachine<Miner>* GetFSM()const{return m_pStateMachine;}

/* EXTRANEIOUS DETAIL OMITTED */
};

```

Notice how the current and global states must be set explicitly when a *StateMachine* is instantiated. The class hierarchy is now like that shown in Figure 2.4.

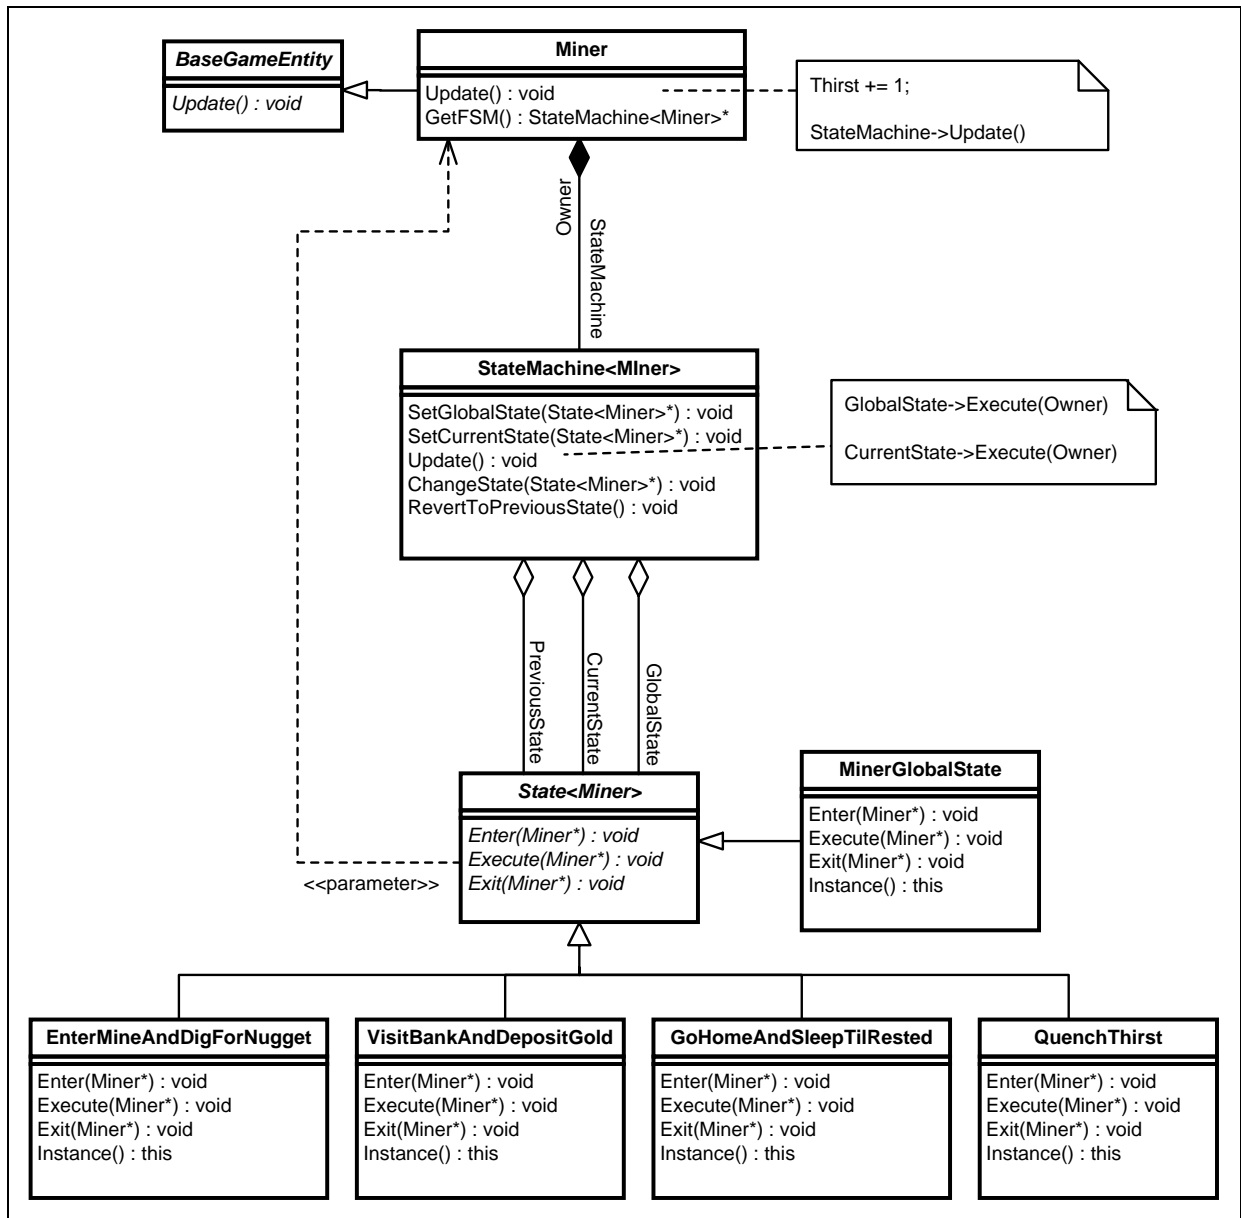


Figure 2.4. The updated design

Introducing Elsa

To demonstrate these improvements, I've created another project: WestWorldWithWoman. In this project, West World has gained another inhabitant, Elsa, the gold miner's wife. Elsa doesn't do much; she's mainly preoccupied with cleaning the shack and emptying her bladder (she drinks way too much cawfee). The state transition diagram for Elsa is shown in Figure 2.5.




Figure 2.5. Elsa's state transition diagram. The global state is not shown in the figure because its logic is effectively implemented in any state and never changed.

When you boot up the project into your IDE, notice how the VisitBathroom state is implemented as a blip state (i.e., it always reverts back to the previous state). Also note that a global state has been defined, WifesGlobalState, which contains the logic required for Elsa's bathroom visits. This logic is contained in a global state because Elsa may feel the call of nature during any state and at any time.

Here is a sample of the output from WestWorldWithWoman.

```

Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Goin' to the bank. Yes siree
Elsa: Walkin' to the can. Need to powda mah pretty li'l nose
Elsa: Ahhhhhh! Sweet relief!
Elsa: Leavin' the john
Miner Bob: Depositin' gold. Total savings now: 4
Miner Bob: Leavin' the bank
Miner Bob: Walkin' to the gold mine
Elsa: Walkin' to the can. Need to powda mah pretty li'l nose
Elsa: Ahhhhhh! Sweet relief!
Elsa: Leavin' the john
Miner Bob: Pickin' up a nugget
Elsa: Moppin' the floor
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Boy, ah sure is thusty! Walkin' to the saloon
Elsa: Moppin' the floor
Miner Bob: That's mighty fine sippin' liquor
Miner Bob: Leavin' the saloon, feelin' good
Miner Bob: Walkin' to the gold mine
  
```

Elsa: Makin' the bed
Miner Bob: Pickin' up a nugget
Miner Bob: Ah'm leavin' the gold mine with mah pockets full o' sweet gold
Miner Bob: Goin' to the bank. Yes siree
Elsa: Walkin' to the can. Need to powda mah pretty li'l nose
Elsa: Ahhhhhh! Sweet relief!
Elsa: Leavin' the john
Miner Bob: Depositin' gold. Total savings now: 5
Miner Bob: Woohoo! Rich enough for now. Back home to mah li'l lady
Miner Bob: Leavin' the bank
Miner Bob: Walkin' home
Elsa: Walkin' to the can. Need to powda mah pretty li'l nose
Elsa: Ahhhhhh! Sweet relief!
Elsa: Leavin' the john
Miner Bob: ZZZZ...

Well, that's it folks. The complexity of the behavior you can create with finite state machines is only limited by your imagination. You don't have to restrict your game agents to just one finite state machine either. Sometimes it may be a good idea to use two FSMs working in parallel: one to control a character's movement and one to control the weapon selection, aiming, and firing, for example. It's even possible to have a state itself contain a state machine. This is known as a hierarchical state machine. For instance, your game agent may have the states Explore, Combat, and Patrol. In turn, the Combat state may own a state machine that manages the states required for combat such as Dodge, ChaseEnemy, and Shoot.